

Contents

1	Evaluation	3
1.1	Evaluation Considerations	4
1.2	Methods	4
1.2.1	Objectives	4
1.2.2	Qualitative	4
1.2.3	Quantitative	5
1.3	Process	6
1.4	Results	6
1.4.1	Qualitative	6
1.4.2	Quantitative	6

Chapter 1

Evaluation

Section ?? described the implementation of the Nickel Language Server addressing the first research question stated in sec. ?. Proving the viability of the result and answering the second research question demands an evaluation of different factors.

Earlier, the most important metrics of interest were identified as:

Usability What is the real-world value of the language server?

Does it improve the experience of developers using Nickel? NLS offers several features, that are intended to help developers using the language. The evaluation should assess whether developers experience any help due to the use of the server.

Does NLS meet its users' expectations in terms of completeness and correctness and behavior? Being marketed as a Language Server, invokes certain expectations due to previous experience with other languages and language servers. Here, the evaluation should show whether NLS lives up to the expectations of its users.

Performance What are the typical latencies of standard tasks? In this context *latency* refers to the time it takes from issuing an LSP command to return of the reply by the server. The JSON-RPC protocol used by the LSP is synchronous, i.e. requires the server to return results of commands in the order it received them. Since most commands are sent implicitly, a quick processing is imperative to avoid commands queuing up.

Can single performance bottlenecks be identified? Single commands with excessive runtimes can slow down the entire communication resulting in bad user experience. Identified issues can guide the future work on the server.

How does the performance of NLS scale for bigger projects? With increasing project sizes the work required to process files increases as well. The evaluation should allow estimates of the sustained performance in real-world scenarios.

Answering the questions above, this chapter consists of two main sections. The first section sec. 1.2 introduces methods employed for the evaluation. In particular, it details the survey (sec. 1.2.2) which was conducted with the intent

to gain qualitative opinions by users, as well as the tracing mechanism (sec. 1.2.3) for factual quantitative insights. Section 1.4 summarises the results of these methods.

1.1 Evaluation Considerations

Different methods to evaluate the abovementioned metrics were considered. While quantifying user experience yields statistically sound insights about the studied subject, it fails to point out specific user needs. Therefore, this work employs a more subjective evaluation based on a standardized experience report focusing on individual features. Contrasting the expectations with experiences allows the implementation more practically and guide the further development by highlighting well executed, immature or missing features.

On the other hand it is more approachable to track runtime performance objectively through time measurements. In fact, runtime behavior was a central assumption underlying the server architecture. As discussed in sec. ?? an eager processing model was chosen over lazy analysis. It was hypothesized that analyzing Nickel source code eagerly allows to perform a single computation ahead of time instead of multiple partial ones. This way both analyzing and querying information could be implemented more efficiently. Moreover, it was assumed that eager computation would have a negligible impact on performance given the relatively small size of Nickel code bases in the current stage.

1.2 Methods

1.2.1 Objectives

The qualitative evaluation was conducted with a strong focus on the first metric in [sec:metrics]. Usability proves hard to quantify, as it is tightly connected to subjective perception, expectations and tolerances. The structure of the survey is guided by two additional objectives, endorsing the separation of individual features. On one hand, the survey should inform the future development of NLS; which feature has to be improved, which bugs exist, what do users expect. This data is important as for NLS both as an LSP implementation for Nickel (affecting the perceived maturity of Nickel) as well as a generic basis for other projects. On the other hand, all features are implemented on top of the same base (cf. sec. ??). The survey should therefore show architectural deficits as well.

The quantitative study in contrast focuses on measurable performance. Similarly to the survey based evaluation it should reveal insight for different features and tasks separately. An additional objective, in line with the definition of the performance metric in [#sec:metrics], is to show the influence of growing file sizes in practice.

1.2.2 Qualitative

Inspired by the work of Leimeister in (**leimeister?**), a survey aims to provide practical insights into the experience of future users. In order to get a clear

picture of the users' needs and expectations independent of the experience, the survey consists of two parts – a pre-evaluation and final survey.

1.2.2.1 Pre-Evaluation

The pre-evaluation introduced participants in brief to the concept of language servers and asked them to write down their understanding of several LSP features. In total, six features were surveyed corresponding to the implementation as outlined in sec. ???. The item for the “Hover” feature for instance reads as follows:

Editors can show some additional information about code under the cursor. The selection, kind, and formatting of that information is left to the Language Server.

What kind of information do you expect to see when hovering code?
Does the position or kind of element matter? If so, how?

Items first introduce a feature on a high level followed by a request to the participant to describe their ideal implementation of the feature.

1.2.2.2 Experience Survey

For the final survey interested participants at Tweag were invited to a workshop introducing Nickel. As a preparation, they were asked to install the LSP. The workshop allowed participants unfamiliar with the Nickel language to use the language and experience NLS. Following the workshop, participants filled in a second survey which focused on particular experiences of every single feature. This evaluation focused on three main aspects. First, the general experience without weighing in expectations. The goal was to assess the extent to which the users were able to use the feature, since all usability metrics as discussed in [sec:metrics] depend on the respective feature being available in the first place. In the same category are the items surveying the perceived performance and stability on a linear scale hinting at possible usability issues. The scales span from “Very slow response” to “Very quick response” and “Never Crashed” to “Always Crashed” respectively. Under the second aspect, the users were asked to explicitly reflect on their expectations in order to contribute to the usability metric. In the final part participants could describe their perceived shortcomings or questions or remarks.

1.2.3 Quantitative

To address the performance metrics introduced in sec. ??, a quantitative study was conducted, that analyzes latencies in the LSP-Server-Client communication. The study complements the subjective reports collected through the survey (cf. sec. 1.2.2.2). The evaluation is possible due to the inclusion of a custom tracing module in NLS. The tracing module is used to create a report for every request, containing the processing time and a measure of the size of the analyzed document. If enabled, NLS records an incoming request with an identifier and time stamp. While processing the request, it adds additional data to the record, i.e., the type of request, the size of the linearization (cf. sec. ??) or processed file and possible errors that occurred during the process. Once the server replies

to a request, it records the total response time and writes the entire record to an external file.

The tracing approach narrows the focus of the performance evaluation to the time spent by NLS. Consequently, the performance evaluation is independent of the LSP client (editor) that is used, which could introduce a bias in the results.

1.3 Process

1.4 Results

1.4.1 Qualitative

1.4.2 Quantitative